

Package ‘HPdregression’

January 20, 2015

Type Package

Title Distributed Regression for Big Data

Version 1.0.0

Date 2015-01-16

Author HP Vertica Analytics Team

Maintainer HP Vertica Analytics Team <distributedRTeam@external.groups.hp.com>

Depends R (>= 3.0.0), distributedR

Description Implementation of distributed generalized linear model. Written using HP Vertica Distributed R package.

License GPL (>= 2) | file LICENSE

R topics documented:

HPdregression-package	1
cv.hpdglm	2
family.hpdglm	3
hpdglm	4
hpdglm.control	6
predict.hpdglm	7
residuals.hpdglm	8
summary.hpdglm	9
v.hpdglm	11
Index	13

HPdregression-package

HPdregression - Distributed Regression for Big Data

Description

HPdregression provides distributed algorithms for regression models. It is written based on the infrastructure created in HP-Lab for distributed computing in R.

Details

Package: HPdregression
 Type: Package
 Version: 1.0.0
 Date: 2015-01-16

Main Functions:

- `hpdglm`: It is a distributed version of `glm`.
- `v.hpdglm`: This function is implemented for evaluating a model built by `hpdglm` using Split-Sample-Validation method.
- `cv.hpdglm`: This function is implemented for evaluating a model built by `hpdglm` using Cross-Validation method.

Author(s)

HP Vertica Analytics Team <distributedRTeam@external.groups.hp.com>

cv.hpdglm

Cross-Validation Method for hpdglm Models

Description

This function is implemented for evaluating a model built by `hpdglm` using Cross-Validation method (This function is only available for complete models). In simple words, the data is randomly divided into K folds, and building model and testing that is repeated K times. Every iteration, one fold is reserved as test data and the rest is used for training. Finally, the prediction costs of the new models on all folds are aggregated.

Usage

```
cv.hpdglm(responses, predictors, hpdglmfit,
          K = 10, sampling_threshold = 1e+06)
```

Arguments

<code>responses</code>	the darray that contains the vector of responses.
<code>predictors</code>	the darray that contains the vector of predictors.
<code>hpdglmfit</code>	a built model of type <code>hpdglm</code> .
<code>K</code>	number of folds in cross validation.
<code>sampling_threshold</code>	threshold for the method of sampling (centralized or distributed). It should be always smaller than <code>1e9</code> . When (<code>blockSize > sampling_threshold</code> <code>nSample > 1e9</code> <code>nSample/K > sampling_threshold</code>), the distributed sampling is selected in which each block is divided to K folds. Here, <code>blockSize</code> is the number of samples in each partition of predictors, and <code>nSample</code> is the total number of samples in predictors.

Details

In order to randomly select the validation set, `sample.int` function of R is used. This function does not support sample space bigger than $1e9$; moreover, it is slow for big numbers. Therefore, when the number of samples in each block of darray becomes bigger than `sampling_threshold`, instead of purely random selection on the master side, each block will randomly contribute its portion to validation set in a distributed fashion. When the ratio of number of samples in each block to the total number of blocks is huge, the skew of randomness is negligible, but performance will be improved.

Value

<code>call</code>	the original call to <code>cv.hpdglm</code>
<code>K</code>	number of folds used at the input
<code>delta</code>	a vector of length two. The first component is the raw validation estimate of prediction error. The second component is the adjusted validation estimate. Lower cost indicates better fitting. In more detail, the 1st-cost is prediction cost of new model on test data, and the 2nd-cost is [1st-cost + (cost of the old model on all data - the cost of the new model on all data)]
<code>seed</code>	the value of <code>.Random.seed</code> when <code>cv.hpdglm</code> was called.

Author(s)

HP Vertica Analytics Team

Examples

```
## Not run:
library(HPdregression)
distributedR_start()
Y <- as.darray(as.matrix(faithful$eruptions),
               c(ceiling(length(faithful$eruption)/4),1))
X <- as.darray(as.matrix(faithful$waiting),
               c(ceiling(length(faithful$waiting)/4),1))

myModel <- hpdglm(Y, X, completeModel=TRUE)
testCV <- cv.hpdglm(Y, X, myModel)

## End(Not run)
```

<code>family.hpdglm</code>	<i>family of an hpdglm model</i>
----------------------------	----------------------------------

Description

Returns the family used for building an hpdglm model.

Usage

```
family.hpdglm(object, ...)
```

Arguments

`object` an hpdglm model

Value

the family of the model

hpdglm	<i>Distributed Generalized Linear Models</i>
--------	--

Description

hpdglm function is intended to be a distributed alternative for glm function.

Usage

```
hpdglm(responses, predictors, family=gaussian, weights=NULL,
na_action="exclude", start=NULL, etastart=NULL, mustart=NULL,
offset=NULL, control=list(...), method="hpdglm.fit.Newton",
completeModel=FALSE, ...)
```

Arguments

responses	the darray that contains the vector of responses.
predictors	the darray that contains the vector of predictors. hpdglm() cannot accept a predictor with constant value. Moreover, a categorical predictor should be decoded (converted to several predictors) before applying hpdglm().
family	it specifies the family function for regression. The supported family-links at the time of this writing are gaussian(identity), binomial(logit), and poisson(log). The mentioned links are the default ones for their families; so, specifying them is optional. The default family is Gaussian.
weights	it is an optional darray of 'prior weights' to be used in the fitting process. It has a single column. The number of rows and its number of blocks should be the same as responses. The values should not be negative (greater than or equal to zero). Weight zero on a sample makes it be ignored.
na_action	it indicates what should happen when the data contain missed values. Values of NA, NaN, and Inf in samples are treated as missed values. There are two options for this argument exclude and fail. When exclude is selected (the default choice), the weight of any sample with missed values will become zero, and that sample will be ignored in the fitting process. In the darray which will be created for residuals, the value corresponding to these samples will be NA. When fail is selected, the function will stop in the case of any missed value in the dataset.
start	starting values for coefficients. It is optional.
etastart	starting values for parameter 'eta' which is used for computing deviance. It should be of type darray. It is optional.
mustart	starting values for mu 'parameter' which is used for computing deviance. It should be of type darray. It is optional.
offset	an optional darray which can be used to specify an <i>a priori</i> known component to be included in the linear predictor during fitting.
control	an optional list of controlling arguments. The optional elements of the list and their default values are: epsilon = 1e-8, maxit = 25, trace = FALSE, rigorous = FALSE.

<code>method</code>	this argument reserved for the future improvement. The only available fitting method at the moment is "hpdglm.fit.Newton". In the future, if we have new developed algorithms, this argument can be used to switch between them.
<code>completeModel</code>	when it is FALSE (default), the output values that have darray structure will be discarded.
<code>...</code>	arguments to be used to form the default <code>control</code> argument if it is not supplied directly.

Details

predictors and responses must align with each other (have the same number of rows and the same number of blocks).

Value

<code>coefficients</code>	calculated coefficients
<code>d.residuals</code>	(only when <code>completeModel=TRUE</code>) the working residuals, that is the residuals in the final iteration of the IWLS fit. Since cases with zero weights are omitted, their working residuals are NA. It is of type darray.
<code>d.fitted.values</code>	(only when <code>completeModel=TRUE</code>) the fitted mean values, obtained by transforming the linear predictors by the inverse of the link function. It is of type darray.
<code>family</code>	the family function used for regression
<code>d.linear.predictors</code>	(only when <code>completeModel=TRUE</code>) the linear fit on link scale. It is of type darray.
<code>deviance</code>	up to a constant, minus twice the maximized log-likelihood.
<code>aic</code>	a version of Akaike's An Information Criterion, minus twice the maximized log-likelihood plus twice the number of parameters, computed by the aic component of the family.
<code>null.deviance</code>	the deviance for the null model, comparable with deviance.
<code>iter</code>	the number of iterations of IWLS used.
<code>prior.weights</code>	(only when <code>completeModel=TRUE</code>) the weights initially supplied. All of its values are 1 if no initial weights used. It is of type darray. The value of weight will become 0 for the samples with invalid data (NA, NaN, Inf).
<code>weights</code>	(only when <code>completeModel=TRUE</code>) the working weights, that is the weights in the final iteration of the IWLS fit. It is of type darray. In order to save memory and execution time, no new darray will be created for weights when the initial weights are all 0 or 1, and it will simply be a reference to <code>prior.weights</code> .
<code>df.residual</code>	the residual degrees of freedom.
<code>df.null</code>	the residual degrees of freedom for the null model.
<code>converged</code>	logical. Was the IWLS algorithm judged to have converged?
<code>boundary</code>	logical. Is the fitted value on the boundary of the attainable values?
<code>responses</code>	(only when <code>completeModel=TRUE</code>) the darray of responses.

predictors	(only when completeModel=TRUE) the darray of predictors.
na_action	this item exists only when a few samples are excluded because of missed data. It is a list containing type "exclude" and the number of excluded samples.
call	the matched call.
offset	(only when completeModel=TRUE) the offset darray used.
control	the value of the control argument used.
method	the name of the fitter function used, currently always "hpdglm.fit.Newton".

Author(s)

HP Vertica Analytics Team

Examples

```
## Not run:
## Example for logistic regression
library(HPdregression)
distributedR_start()
# selecting 4 partitions
Y <- as.darray(as.matrix(mtcars$am), c(ceiling(length(mtcars$am)/4), 1))
colnames(Y) <- "am"
X <- as.darray(as.matrix(cbind(mtcars$wt, mtcars$hp)),
               c(ceiling(length(mtcars$hp)/4), 2))
colnames(X) <- c("wt", "hp")
# building logistic regression model
myModel <- hpdglm(Y, X, binomial)
summary(myModel)

## Example for linear regression
require(MASS)
y <- matrix(Boston$medv, ncol=1)
n <- nrow(y)
x <- matrix(c(Boston$rad, Boston$crim, Boston$ptratio, Boston$dis), ncol=4)
# selecting the number of partitions the same as the number of R-executors
nparts <- sum(distributedR_status()$Inst)
block_row <- ceiling(n/nparts)
# creating the darray of response
Y <- as.darray(y, c(block_row, 1))
colnames(Y) <- c("medv")
# creating the darray of predictors
X <- as.darray(x, c(block_row, 4))
colnames(X) <- c("rad", "crim", "ptratio", "dis")
# building linear regression model
reg <- hpdglm(Y, X)
summary(reg)

## End(Not run)
```

Description

Auxiliary function for `hpdglm` fitting. Typically only used internally by `hpdglm.fit`, but may be used to construct a `control` argument to either function.

Usage

```
hpdglm.control(epsilon = 1e-08, maxit = 25, trace = FALSE,
               rigorous = FALSE)
```

Arguments

<code>epsilon</code>	It is used to adjust desired accuracy of the result.
<code>maxit</code>	It is the maximum, number of iterations before achieving the desired accuracy.
<code>trace</code>	When this argument is true, intermediate steps of the progress are displayed.
<code>rigorous</code>	When this argument is true, some extra checks are performed during fitting procedure. For example, μ and η may be validating in each iteration to check if the fitted values are outside of the domain. Usually these checks are time consuming; therefore, the default value for this argument is FALSE.

Value

A list with components named as the arguments.

Examples

```
## Not run:
library(HPdregression)
distributedR_start()
Y <- as.darray(as.matrix(mtcars$am), c(ceiling(length(mtcars$am)/4), 1))
X <- as.darray(as.matrix(cbind(mtcars$wt, mtcars$hp)),
              c(ceiling(length(mtcars$hp)/4), 2))

myModel <- hpdglm(Y, X, binomial, control=list(epsilon=1e-02, maxit=5,
                                              trace=FALSE, rigorous=TRUE))

## End(Not run)
```

predict.hpdglm	<i>Predict Method for hpdglm fits</i>
----------------	---------------------------------------

Description

It produces predicted values, obtained by evaluating the regression function on provided new data.

Usage

```
predict.hpdglm(object, newdata, type = c("link", "response"),
               na.action = na.pass, mask = NULL, trace = TRUE, ...)
```

Arguments

<code>object</code>	a built model of type <code>hpdglm</code> .
<code>newdata</code>	a matrix or a darray containing predictors of new samples.
<code>type</code>	the type of prediction required which can be "link" or "response".
<code>na.action</code>	a function to determine what should be done with missing values. At this version it is always <code>na.pass</code> (reserved for future improvement).
<code>mask</code>	a darray with a single column, and 0 or 1 as the value of its elements. It indicates which samples (rows) should be considered in the calculation.
<code>trace</code>	when this argument is true, intermediate steps of the progress are displayed.
<code>...</code>	further arguments passed to or from other methods.

Details

This function produces predicted values, obtained by evaluating the regression function on provided new data. New data can be either a darray or a normal matrix.

Value

The output is a matrix or a darray, depending to the type of `newdata`, which contains predicted values for response.

Author(s)

HP Vertica Analytics Team

Examples

```
## Not run:
library(HPdregression)
distributedR_start()
Y <- as.darray(as.matrix(faithful$eruptions),
               c(ceiling(length(faithful$eruption)/4),1))
X <- as.darray (as.matrix(faithful$waiting),
               c(ceiling(length(faithful$waiting)/4),1))

myModel <- hpdglm(Y, X)
newSamples <- matrix(c(1:3),,1)
predict(myModel, newSamples, "link")

## End(Not run)
```

`residuals.hpdglm` *Extract Residuals of an hpdglm Model*

Description

This function extracts model residuals of an `hpdglm` model in a darray. The abbreviated function is `resid`. This function is only available for complete models.

Usage

```
residuals.hpdglm(object, type = c("deviance", "pearson",
                                   "working", "response", "partial"), trace=FALSE, ...)
```

Arguments

object	an hpdglm model
type	can be "deviance", "pearson", "working", "response", or "partial".
trace	when TRUE, intermediate steps of the progress are displayed.

Value

darray of residuals

Examples

```
## Not run:
library(HPdregression)
distributedR_start()
Y <- as.darray(as.matrix(mtcars$am),
               c(ceiling(length(mtcars$am)/4), 1))
X <- as.darray(as.matrix(cbind(mtcars$wt, mtcars$hp)),
               c(ceiling(length(mtcars$hp)/4), 2))

myModel <- hpdglm(responses=Y, predictors=X,
                  family=binomial(logit), completeModel=TRUE)
res <- resid(myModel)

## End(Not run)
```

summary.hpdglm

*Summarizing The Model Made by hpdglm***Description**

This function prints a summary of the learned model.

Usage

```
summary.hpdglm(object, dispersion = NULL, correlation = FALSE,
               symbolic.cor = FALSE, trace=FALSE, ...)
```

Arguments

object	a result of a call to hpdglm .
dispersion	the dispersion parameter for the family used. Either a single numerical value or NULL (the default), when it is inferred from object).
correlation	logical; if TRUE, the correlation matrix of the estimated parameters is returned and printed.
symbolic.cor	logical. If TRUE, print the correlations in a symbolic form (see symnum) rather than as numbers.
trace	when TRUE, intermediate steps of the progress are displayed.

Value

summary.hpdglm returns an object of class "summary.hpdglm", a list with components

call	the component from object.
family	the component from object.
deviance	the component from object.
df.residual	the component from object.
null.deviance	the component from object.
df.null	the component from object.
deviance.resid	(only when completeModel=TRUE) the deviance residuals: see residuals.hpdglm .
coefficients	the matrix of coefficients, standard errors, z-values and p-values. Aliased coefficients are omitted.
aliased	named logical vector showing if the original coefficients are aliased.
dispersion	either the supplied argument or the inferred/estimated dispersion if the latter is NULL.
df	a 3-vector of the rank of the model and the number of residual degrees of freedom, plus number of non-aliased coefficients.
cov.unscaled	the unscaled (dispersion = 1) estimated covariance matrix of the estimated coefficients.
cov.scaled	ditto, scaled by dispersion.
correlation	(only if correlation is true.) The estimated correlations of the estimated coefficients.
symbolic.cor	(only if correlation is true.) The value of the argument symbolic.cor.
minMax	it contains the minimum and the maximum of the deviance residuals.

Author(s)

HP Vertica Analytics Team

Examples

```
## Not run:
require(distributedR)
distributedR_start()
Y <- as.dobject(as.matrix(faithful$eruptions),
                c(ceiling(length(faithful$eruption)/4),1))
X <- as.dobject(as.matrix(faithful$waiting),
                c(ceiling(length(faithful$waiting)/4),1))

myModel <- hpdglm(responses=Y, predictors=X)
summary(myModel)

## End(Not run)
```

Description

This function is implemented for evaluating a model built by hpdglm using Split-Sample-Validation method (This function is only available for complete models). In simple words, a percent of the data is randomly selected as test data, and a new model is built using the rest. Finally, the prediction cost of the new model on the test data is measured.

Usage

```
v.hpdglm(responses, predictors, hpdglmfit,
          percent = 30, sampling_threshold = 1e+06)
```

Arguments

responses	the darray that contains the vector of responses.
predictors	the darray that contains the vector of predictors.
hpdglmfit	a built model of type hpdglm.
percent	the percent of data which should be set aside for validation.
sampling_threshold	threshold for the method of sampling (centralized or distributed). It should be always smaller than 1e9. When (blockSize > sampling_threshold nSample > 1e9), the distributed sampling is selected in which a percent of each block is selected for test data. Here, blockSize is the number of samples in each partition of predictors, and nSample is the total number of samples in predictors.

Details

In order to randomly select the validation set, sample.int function of R is used. This function does not support sample space bigger than 1e9; moreover, it is slow for big numbers. Therefore, when the number of samples in each block of darray becomes bigger than sampling_threshold, instead of purely random selection on the master side, each block will randomly contribute its portion to validation set in a distributed fashion. When the ratio of number of samples in each block to the total number of blocks is huge, the skew of randomness is negligible, but performance will be improved.

Value

call	the original call to v.hpdglm
percent	the percent value used at the input
delta	a vector of length two. The first component is the raw validation estimate of prediction error. The second component is the adjusted validation estimate. Lower cost indicates better fitting. In more detail, the 1st-cost is prediction cost of new model on test data, and the 2nd-cost is [1st-cost + (cost of the old model on all data - the cost of the new model on all data)]
seed	the value of .Random.seed when v.hpdglm was called.

Author(s)

HP Vertica Analytics Team

Examples

```
## Not run:
library(HPdregression)
distributedR_start()
Y <- as.darray(as.matrix(faithful$eruptions),
               c(ceiling(length(faithful$eruption)/4),1))
X <- as.darray(as.matrix(faithful$waiting),
               c(ceiling(length(faithful$waiting)/4),1))

myModel <- hpdglm(Y, X, completeModel=TRUE)
testV <- v.hpdglm(Y, X, myModel)

## End(Not run)
```

Index

*Topic **Cross-Validation**

`cv.hpdglm`, 2

*Topic **Distributed R**

`cv.hpdglm`, 2

`hpdglm`, 4

`hpdglm.control`, 6

`residuals.hpdglm`, 8

`v.hpdglm`, 11

*Topic **Distributed regression**

`summary.hpdglm`, 9

*Topic **Generalized Linear Regression**

`hpdglm`, 4

*Topic **Split-Sample-Validation**

`v.hpdglm`, 11

*Topic **distributed R, distributed Regression, Big Data Analytics**

`HPdregression-package`, 1

*Topic **distributed R**

`predict.hpdglm`, 7

*Topic **family**

`family.hpdglm`, 3

*Topic **hpdglm models**

`summary.hpdglm`, 9

*Topic **hpdglm model**

`cv.hpdglm`, 2

`hpdglm.control`, 6

`residuals.hpdglm`, 8

`v.hpdglm`, 11

*Topic **hpdglm**

`family.hpdglm`, 3

*Topic **models**

`predict.hpdglm`, 7

*Topic **regression**

`predict.hpdglm`, 7

`cv.hpdglm`, 2

`deviance.hpdglm(hpdglm)`, 4

`family.hpdglm`, 3

`fitted.hpdglm(hpdglm)`, 4

`hpdglm`, 4, 7, 9

`hpdglm.control`, 6

`hpdglm.fit`, 7

`HPdregression`

(*HPdregression-package*), 1

`HPdregression-package`, 1

`predict.hpdglm`, 7

`print.hpdglm(hpdglm)`, 4

`print.summary.hpdglm`

(*summary.hpdglm*), 9

`resid(residuals.hpdglm)`, 8

`residuals(residuals.hpdglm)`, 8

`residuals.hpdglm`, 8, 10

`summary.hpdglm`, 9

`symnum`, 9

`v.hpdglm`, 11

`weights.hpdglm(hpdglm)`, 4